

RAG System Test Document

This is a test document for our RAG (Retrieval Augmented Generation) system.

Page 1: Introduction

RAG combines retrieval and generation. It first retrieves relevant documents from a knowledge base, then uses a language model to generate answers based on the retrieved context.

Key benefits of RAG:

1. Reduces hallucinations
2. Provides source citations
3. Easy to update knowledge
4. Cost-effective vs fine-tuning

Our system uses:

- BGE-M3 for embeddings (1024 dimensions)
- PostgreSQL with pgvector for storage
- Ollama with Llama 3.1 for generation
- FastAPI for the API layer

This document tests the full upload pipeline including text extraction, chunking, embedding generation, and vector storage.